

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/312616139>

# Kokkos, Modern c++ and performance portability

Presentation · January 2017

DOI: 10.13140/RG.2.2.26312.11523

---

CITATIONS

0

---

READS

732

1 author:



[Pierre Kestener](#)

Atomic Energy and Alternative Energies Commission

71 PUBLICATIONS 793 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:



Performance portable implementation for high-order numerical scheme. [View project](#)



CanoP AMR [View project](#)

# Kokkos, Modern C++, performance portability, ...

Pierre Kestener<sup>1</sup>

<sup>1</sup>CEA Saclay, DSM, Maison de la Simulation

PATC, January, 16-18th, 2017



## Schedule

- **Monday + Tuesday morning:** NVidia OpenACC tutorial
- **Tuesday afternoon + Wednesday:** Kokkos tutorial
  - Introduction performance portability
  - IBM Power8 + Nvidia Pascal P100 platform : short overview
  - **Kokkos: features overview**
  - **Hands-on 0: retrieve Kokkos sources**, how to build, how to run a *helloworld* application, explore different configurations
  - **Hands-on 1:** cross-checking **Kokkos + hwloc** is OK
  - **Replay some tutorial slides from SC2016 for deeper Kokkos concepts**
  - **Hands-On 2:** Simple example **SAXPY**  
⇒ simplest computing kernel in Kokkos
  - **Hands-On 3:** Simple example **Mandelbrot set**  
⇒ 1D Kokkos::View + linearized index (+ asynchronous execution)
  - **Hands-On 4:** Simple examples **Stencil + Finite Difference**  
⇒ 2D Kokkos::View
  - **Hands-On 5: Laplace exercise**  
⇒ pure Kokkos versus Kokkos + MPI + hwloc (multiGPU)
  - **Hands-On 6:** CSCS miniApp: **Fisher equation solver**  
⇒ use Kokkos lambda
  - **Hands-On 7:** CFD miniApp: **Euler solver**  
⇒ performance measurement for several Kokkos backends (OpenMP, CUDA)



## IBM Power8 / Nvidia Pascal P100

- **Kokkos training material archive (last up-to-date version) is on ouessant:**  
[/pwrwork/workshops/patc-201701/kokkos/training\\_kokkos.tar.gz](/pwrwork/workshops/patc-201701/kokkos/training_kokkos.tar.gz)
- Use material from IBM/NVidia<sup>1</sup>, gives detail on the platform  
See file: <doc/ouessant/Introduction.pdf> in archive
- Minimal information about software environment, how to build and run an application, submit a job on machine ouessant  
See file:  
[doc/ouessant/Ouessant-Application\\_User\\_Guide-16-12-1.pdf](doc/ouessant/Ouessant-Application_User_Guide-16-12-1.pdf)


---

<sup>1</sup>Thanks to Nicolas Tallet (IBM)




## Kokkos: a programming model for performance portability

- **Kokkos** is a C++ **library** with **parallel algorithmic patterns** AND **data containers** for **node-level parallelism**.
- Implementation relies heavily on **meta-programing** to derive native **low-level code** (OpenMP, Pthreads, CUDA, ...) and adapt data structure **memory layout** at compile-time
- Core developers at **SANDIA NL** (H.C. Edwards, C. Trott)



# Kokkos: a programming model for performance portability

- **Open source**, <https://github.com/kokkos/kokkos>
- Primarily developed as a base building layer for **generic high-performance parallel linear algebra** in [Trilinos](#)
- Also used in molecular dynamics code, e.g. [LAMMPS](#)
- Goal: **ISO/C++ 2020 Standard** subsumes Kokkos abstractions

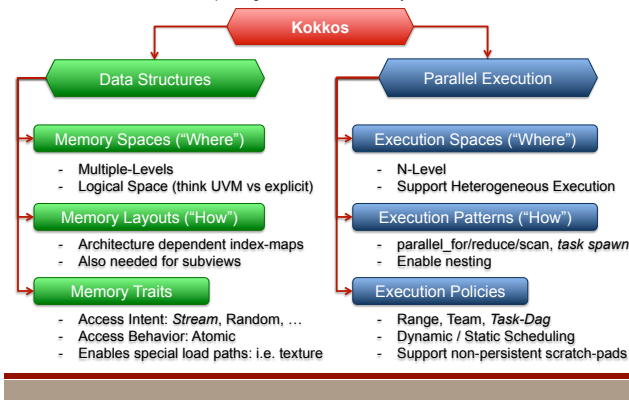


# Kokkos: a programming model for performance portability

## Performance Portability through Abstraction



Separating of Concerns for Future Systems...



reference:

<https://cfwebprod.sandia.gov/cfdocs/CompResearch/docs/Kokkos-Multi-CoE.pdf>



## Kokkos Concepts (1) - the abstract machine model

- Kokkos defines an abstract machine model for future large shared-memory nodes made of
  - **latency-oriented cores** (contemporary CPU core)
  - **throughput-oriented cores** (GPU, ...)




Figure: Conceptual model of a future HPC node. (Kokkos User's Guide).



## Kokkos Concepts (2) - What is a device ?

- A **Kokkos device**:
- From a C++ API design point of view, Kokkos defines several c++ class for a **device** in `core/src`, e.g.
  - `Kokkos::Cuda`, `Kokkos::OpenMP`, `Kokkos::Pthreads`, `Kokkos::Serial`
  - *device = execution space + memory space*
- Each *Kokkos device* pre-defines some types
- Example **Kokkos device** (not required for a user, only Kokkos developer), e.g.

```
class Cuda {  
public:  
    // Tag this class as a kokkos execution space  
    typedef Cuda          execution_space ;  
  
    #if defined( KOKKOS_USE_CUDA_UVM )  
        // This execution space's preferred memory space.  
        typedef CudaUVMSpace    memory_space ;  
    #else  
        // This execution space's preferred memory space.  
        typedef CudaSpace       memory_space ;  
    #endif  
  
    // This execution space preferred device_type  
    typedef Kokkos::Device<execution_space,memory_space> device_type;  
  
    // The size_type best suited for this execution space.  
    typedef memory_space::size_type    size_type ;  
  
    // This execution space's preferred array layout.  
    typedef LayoutLeft                array_layout ;  
    ...  
} // end class Cuda
```



## Kokkos Concepts (3) - execution space, memory space

- **Execution space:** Where should a parallel construct (`parallel_for`, `parallel_reduce`, ...) be executed
  - Special case: `class HostSpace`, special device (always defined) where execution space is either (Serial, Pthread or OpenMP).
  - Each execution space is equipped with a fence: `Kokkos::Cuda::fence()`
- **Memory space:** Where / how data are allocated in memory (`HostSpace`, `CudaSpace`, `CudaUVMSpace`, `CudaHostPinnedSpace`, `HBWSpace`, ...)
- **Memory layout** (come back later on that)
- Other concepts:
  - Execution policy: used to modify a parallel thread dispatch
- **Multiple execution / memory space** can be used in a single application  
See for example in Kokkos sources  
`example/tutorial/Advanced_View/07_Overlapping_DeepCopy`  
Though, take care that currently, Cuda stream are not completely mapped into Kokkos API<sup>2</sup>; meanwhile Cuda streams can be used directly (but loses portability);

---

<sup>2</sup>Will be implemented in the coming months



## Hands-On 0: Build kokkos

- **Kokkos is still experimental, but moving fast**
- **1. Get Kokkos sources, development branch**
  - **Practicals on ouessant:**
    1. `mkdir $HOME/kokkos-tutorial; cd $HOME/kokkos-tutorial`  
some kokkos tutorial examples have a Makefile configured for using that precise location.
    2. `git clone https://github.com/kokkos/kokkos`
    3. `cd kokkos; git checkout develop`
- **2. Build configuration**
  - About build system, several ways to use/build Kokkos
    - ① **CMake:** only when Kokkos is build inside [Trilinos](#),
    - ② **Regular standalone utilization (installed Kokkos):** use `generate_makefile.bash`, then `make kokkoslib; make install`  
Then use a *modulefile* to configure the environment
    - ③ **Embedded Kokkos source files in your application** - mostly usefull in tutorial.
  - We will use 2. and 3.



## Hands-On 0: Build kokkos (2)

### About Makefile variable for building on multiple architectures

- The following variables are useful when building some of the tutorial examples :
  - KOKKOS\_PATH: path to Kokkos source dir
  - KOKKOS\_DEVICES: define possible execution spaces: CUDA, OpenMP, Pthreads, Serial, ...
  - KOKKOS\_ARCH: used to customize compiler flags; e.g. Power8, Kepler35, SNB, KNL, ARMv80, ...
- When building for **CUDA device**, you'll need to use Kokkos' own compiler wrapper: **nvcc\_wrapper** (included in Kokkos sources)
- **When building Kokkos and aiming at an installed Kokkos**, the same information (in a different form) is passed to script `generate_makefile.bash`  
Just type `./generate_makefile.bash -help` at top-level Kokkos sources
- **When using Kokkos embedded in your application**, these variables must be set on the make command line.



## Hands-On 0: Build kokkos (3)

### • Example build configurations (for an installed Kokkos)

- For ouessant, see file `doc/readme_build_kokkos_ouessant` in the provided archive
- Serial (mostly for testing)  
`../generate_makefile.bash -with-serial`  
`-prefix=$HOME/local/kokkos_serial`
- **OpenMP**  
`../generate_makefile.bash -with-openmp`  
`-prefix=$HOME/local/kokkos_openmp_dev`
- **CUDA (+ OpenMP)**; typical configuration  
`../generate_makefile.bash -with-cuda -arch=Pascal60`  
`-prefix=$HOME/local/kokkos_cuda_lambda_openmp`  
`-with-cuda-options=enable_lambda -with-openmp -with-hwloc=/usr`

- **After installation** (`make kokkoslib`; `make install`;) the file `Makefile.kokkos` is created, and designed to be reused in your application build system.

### • 2 choices for integrating Kokkos in your app:

- Use an existing Makefile from Kokkos tutorial, examples, ...
- Use your own build system: there can be a quite large combinatorics of DEVICES, ARCH, compilers, compiler options, ...



## Kokkos - Documentation

- PDF documentation in kokkos source tree : `doc/Kokkos_PG.pdf` (programming guide)
- Doxygen can only be built from inside Trilinos source tree  
Version of the day can be browsed at  
<https://trilinos.org/docs/dev/packages/kokkos/doc/html/index.html>
- Kokkos source code itself, reading unit tests code is also helpful

Additionnal resources:

- Tutorial slides and codes:  
<https://github.com/kokkos/kokkos-tutorials>



## Kokkos - initialize / finalize

- Kokkos::initialize / finalize

```
#include <Kokkos_Macros.hpp>
#include <Kokkos_Core.hpp>

int main(int argc, char* argv[]) {
    // default: initialize the host exec space
    // What exactly gets initialized depends on how kokkos
    // was built, i.e. which options was passed to
    // generate_makefile.bash
    Kokkos::initialize();
    ...
    Kokkos::finalize();
}
```

- **What's happening inside Kokkos::initialize**

- Defines **Default Device** / **DefaultExecutionSpace** **Default memory space** (as specified when kokkos itself was built, by order of priority: Cuda > OpenMP > Pthreads > Serial)  
e.g. if `--with-cuda` was not pass to `generate_makefile.bash`, but `--with-openmp` was, then **DefaultExecutionSpace** is OpenMP
- You can activate several execution spaces (recommended)
- all this information provided at compile time will internally be used inside Kokkos sources as default (hidden) template parameters



## Kokkos - initialize / finalize

- Kokkos::initialize / finalize (most of the time OK)

```
#include <Kokkos_Macros.hpp>
```

```
#include <Kokkos_Core.hpp>
```

```
int main(int argc, char* argv[]) {  
    // default: initialize the host exec space  
    // What exactly gets initialized depends on how kokkos  
    // was built, i.e. which options was passed to  
    // generate_makefile.bash  
    Kokkos::initialize();  
    ...  
    Kokkos::finalize();  
}
```

- **Fine control of initialization:**

- Kokkos::initialize(argc, argv);  
User can change/fix e.g. number OpenMP threads on the application's command line
- This is regular initialization. If available **hwloc** is used to provide default hardware locality:
  - For OpenMP exec space: number of threads (default is all CPU cores)  
NB: usual environment variables (e.g. OMP\_NUM\_THREADS, GOMP\_CPU\_AFFINITY can (of course) also be used
  - Mapping between GPUs and MPI task





## Kokkos - initialize / finalize

- **Advanced initialization** with **OpenMP + CUDA**

Needed/usefull to be able to execution computation on both HOST / GPU

```
#if defined( KOKKOS_HAVE_CUDA )  
Kokkos::HostSpace::execution_space::initialize(teams*num_threads);  
Kokkos::Cuda::SelectDevice select_device(device);  
Kokkos::Cuda::initialize(select_device);  
#elif defined( KOKKOS_HAVE_OPENMP )  
Kokkos::OpenMP::initialize(teams*num_threads);  
#elif defined( KOKKOS_HAVE_PTHREAD )  
Kokkos::Threads::initialize(teams*num_threads);  
#endif
```



## Kokkos - initialize / finalize with MPI

- **Advanced initialization** with **MPI + Kokkos/CUDA version 1 : implicit mapping**

Don't do anything special, let Kokkos through hwloc chose the GPU

```
// Just checking how Kokkos+hwloc performed
// the MPI rank - GPU mapping
int cudaDeviceId;
cudaGetDevice(&cudaDeviceId);
std::cout << "I'm MPI task #" << rank << " pinned to GPU #" << cudaDeviceId << "\n";
```

- **Advanced initialization** with **MPI + Kokkos/CUDA version 2 : explicit mapping** (we will come back into that with example code)

```
// MPI initialized above

// probe the number of CUDA device (i.e. GPUs)
const int ngpu = Kokkos::Cuda::detect_device_count();

// provide a mapping 1 MPI task <-> 1 GPU
const int cuda_device_rank = pre_mpi_local_rank % ngpu ;

// each MPI task initialize the selected device id
Kokkos::Cuda::initialize(
Kokkos::Cuda::SelectDevice( cuda_device_rank ) );
```

- In any case, cross-check this information with the job scheduler, e.g. mpirun --report-bindings



## Hands-On 1 : query\_device

**Purpose:** just cross-checking Kokkos/Hwloc is working OK

- We will first re-use material from Kokkos github repository.
- On your home, on ouessant:
  - 1 `mkdir kokkos-tutorial; cd kokkos-tutorial`
  - 2 `git clone https://github.com/kokkos/kokkos.git`  
# **Don't try to build kokkos here (for now)**



## Hands-On 1 : query\_device

### Purpose: just cross-checking Kokkos/Hwloc is working OK

- Kokkos sources will be built by the application Makefile
- `cd $HOME/kokkos-tutorial/kokkos/example/query_device`
- `open query_device.cpp`; no computations, it just prints hardware information
- ❶ **Default serial build (with hwloc):** `make KOKKOS_USE_TPLS="hwloc"`  
How many NUMA / Cores / Hyperthreads on power8 CPU ?  
What is the current SMT mode on a ouessant login node ? (use command `ppc64_cpu --smt` or `ppc64_cpu -info`)
- ❷ **OpenMP build (with hwloc):** `make KOKKOS_USE_TPLS="hwloc"`  
`KOKKOS_DEVICES=OpenMP` (off course, exact same information obtained)
- ❸ **CUDA/OpenMP build (with hwloc):** `make KOKKOS_USE_TPLS="hwloc"`  
`KOKKOS_DEVICES=Cuda,OpenMP`; rerun and you should get information about the CPU+GPU configuration
- **Take some time to have a look at Makefile.**  
Note that latter when using an installed kokkos library, we won't need to set architecture or device related variables on the command line .



## Hands-On 1 : query\_device

**Purpose:** just cross-checking Kokkos/Hwloc is working OK

- **What happens if hwloc is not activated ?**
- Edit file `query_device.cpp` and do the following modification:

- 1 Add `Kokkos::initialize(argc, argv);` after `MPI_Init`
- 2 Add `Kokkos::finalize();` before `MPI_Finalize`
- 3 change

```
#if defined( KOKKOS_HAVE_CUDA )  
    Kokkos::Cuda::print_configuration( msg );  
#else  
    Kokkos::OpenMP::print_configuration( msg );  
#endif
```

- Rebuild 1 **without HWLOC**: make `KOKKOS_DEVICES=OpenMP`

```
Kokkos::OpenMP KOKKOS_HAVE_OPENMP thread_pool_topology[ 1 x 80 x 1
```

- Rebuild 2 **with HWLOC**: make `KOKKOS_DEVICES=OpenMP`  
`KOKKOS_USE_TPLS="hwloc"`

```
hwloc( NUMA[2] x CORE[10] x HT[4] )
```

```
Kokkos::OpenMP KOKKOS_HAVE_OPENMP hwloc[2x10x4] hwloc_binding_enabl
```

- As already said: processor affinity is crucial to performance



## Kokkos data Container (1)

**Kokkos::View<...>** is **multidimensionnal data container** with **hardware adapted memory layout**

- `Kokkos::View<double **> data("data",NX,NY);` : 2D array with sizes known at runtime
- `Kokkos::View<double *[3]> data("data",NX);` : 2D array with first size known at runtime ( $NX$ ), and second known at compile time (3).
- How do I access data ?  $data(i, j)$  !
- **Which memory space ?** By default, the default device memory space !  
Want to enforce in which memory space lives the view ? `Kokkos::View<..., Device>`: if a second template parameter is given, Kokkos expects a Device (e.g. `Kokkos::OpenMP`, `Kokkos::Cuda`, ...)
- `Kokkos::View` are **small**, designed as reference to allocated memory buffer
  - **View = pointer to data + array shape**
  - assignment is fast (shallow copy + increment ref counter)<sup>3</sup>
- `Kokkos::View` are designed to be pass by value to a function.

---

<sup>3</sup>NB: same behaviour as in python for example



## Kokkos data Container (2)

- Concept of **memory layout**:
- **Memory layout is crucial for performance**:
  - **LayoutLeft**:  $data(i, j, k)$  uses linearized index as  $i + NX * j + NX * NY * k$  (column-major order)
  - **LayoutRight**:  $data(i, j, k)$  uses linearized index as  $k + NZ * j + NZ * NY * i$  (row-major order)
- **Kokkos::View<int\*\*, Kokkos::OpenMP>** defaults with **LayoutRight**; a single thread access contiguous entries of the array. Better for cache and avoid sharing cache lines between threads.
- **Kokkos::View<int\*\*, Kokkos::Cuda>** defaults **LayoutLeft** so that consecutive threads in the same warp access consecutive entries in memory; try to ensure memory coalescence constraint
- You can if you like, still enforce memory layout yourself (or just use 1D Views, and compute index yourself);  
We will see the 2 possibilities with the miniApp on the Fisher equation



## Kokkos data Container (3)

- `Kokkos::View<...>` are reference-counted
- By default do a **shallow copy**

```
Kokkos::View<int*>("a",10);  
Kokkos::View<int*>("b",10);  
a = b; // a now points to b (ref counter incremented by 1)
```

- **Deep copy** must be explicit:

```
Kokkos::deep_copy(dest,src);
```

- **Usefull when copying data from one memory space to another**  
e.g. **from HostSpace to CudaSpace**
- When `dest` and `src` are in the same memory space, it does nothing! (usefull for portability, see example in miniapps later)





## Kokkos data Container (4)

- A verbose **Kokkos::View** declaration example:

```
Kokkos::View<double*, Kokkos::LayoutLeft, Kokkos::CudaSpace> a;
```

- **What ?** a data type
- **How ?** a memory layout
- **Where ?** a memory space
- the last two template parameters are optionnal (have default values)
- There is actually a 4th template parameter for Memory traits (e.g. atomic access)
- **Kokkos::DualView<...>**: usefull when porting an application incrementally, adata container on two different memory space.  
see `tutorial/Advanced_Views/04_dualviews/dual_view.cpp`
- **Kokkos::UnorderedMap<...>**
- Can also define **subview (array slicing, no deep copy)**. See exercice about Mandelbrot set.



## Kokkos data Container (5)

- **What types of data may a View contain ?**

C++ Plain Old Data (POD), i.e. basically compatible with C language:

- Can be allocated with `std::malloc`
- Can be copied with `std::memmove`
- POD in C++11:
  - a trivial type (no virtual member functions, no virtual base class)
  - a standard layout type
- C++11: How to check if a given class A is POD ?

```
#include <type_traits>
```

```
class A { ... }
```

```
std::cout << "is class A POD ? " << std::is_pod<A>::value << "\n";
```



## Kokkos data Container (6)

### Interoperability

- **With a legacy API** `void legacyFunction(int * data, int size)`  
how to retrieve a raw pointer from a `Kokkos::View<int *> data`:  
`int *raw_ptr = data.ptr_on_device()`  
This is not recommended. No more reference counting.



## Kokkos compute Kernels - parallel dispatch (1)

- **3 types of parallel dispatch**

- `Kokkos::parallel_for`
- `Kokkos::parallel_reduce`
- `Kokkos::parallel_scan`

- A dispatch needs as input

- **an execution policy:** e.g. a range (can simply be an integer), team of threads, ...
- **a body:** specified as a lambda function or a functor

- Very important: launching a kernel (thread dispatching) is by default asynchronous



## Kokkos compute Kernels - parallel dispatch (2)

### How to specify a compute kernel in Kokkos ?

#### 1 Use Lambda functions.

NB: a lambda in c++11 is an unnamed function object capable of capturing variables in scope.

```
Kokkos::parallel_for (100, KOKKOS_LAMBDA (const int i) {  
    data(i) = 2*i;  
});
```

**Here we do 2 things in 1 step: define the computation body (lambda func) and launch computation.**

#### 2 Use a C++ functor class.

A functor is a class containing a function to execute in parallel.

```
class FunctorType {  
public:  
    KOKKOS_INLINE_FUNCTION  
    void operator() ( const int i ) const ;  
};  
...  
FunctorType func;  
Kokkos::parallel_for (100, func);
```

**Note: 100 here is the simplest way to specify an execution policy**



## Kokkos compute Kernels - parallel dispatch (3)

### Notes on macros defined in `core/src/Kokkos_Macros.hpp`

- `KOKKOS_LAMBDA` is a macro which provides a compiler-portable way of specifying a lambda function with **capture-by-value closure**.
  - `KOKKOS_LAMBDA` must be used at the most outer parallel loop; inside a lambda one can call another lambda
- `KOKKOS_INLINE_FUNCTION` `void operator() (...) const;`  
this macro helps providing the necessary compiler specific *decorators*, e.g. `__device__` for Cuda to make sure the body can be turns into a Cuda kernel.
  - macro `KOKKOS_INLINE_FUNCTION` must be applied to any function call inside a parallel loop



## Kokkos compute Kernels - parallel dispatch (4)

### Lambda or Functor: which one to use in Kokkos ? Both !

#### 1 Use Lambda functions.

- easy way for small compute kernels
- For GPU, requires Cuda 7.5 (8.0 is current and latest CUDA version)

#### 2 Use a C++ functor class.

- More flexible, allow to design more complex kernel



## Kokkos compute Kernels - parallel dispatch (5)

### About Kokkos::parallel\_reduce with lambda

- As for `parallel_for`, loop body can be specified as a lambda, or a functor;  
here is the lambda way when reduce operation is sum:

```
Kokkos::parallel_reduce (100, KOKKOS_LAMBDA (const int i, int &local_sum)  
    local_sum += data(i);  
}, sum);
```





## Kokkos compute Kernels - parallel dispatch (6)

### About Kokkos::parallel\_reduce functor

- Kokkos supplies a default `init / join` operator which is `operator+`
- If the reduce operator is not trivial (i.e. not a sum)  $\Rightarrow$  you need to define methods `init` and `join`

```
class ReduceFunctor {  
    public:  
        // declare a constructor ...  
        KOKKOS_INLINE_FUNCTION void  
        operator() (const int i, data_t &update) const {...}  
  
        // How to join/combine intermediate reduce from different threads  
        KOKKOS_INLINE_FUNCTION void  
        join(volatile data_t &dst, const volatile data_t &src) const {...}  
  
        // how each thread initializes its reduce result  
        KOKKOS_INLINE_FUNCTION void  
        init(const volatile data_t &dst) const {...}  
}
```



## Kokkos compute Kernels - parallel dispatch (7)

### Parallel dispatch - execution policy

- Remember that an execution policy specifies **how** a parallel dispatch is done by the device
- Range policy:** from...to  
no prescription of order of execution nor concurrency; allows to adapt to the actual hardware; e.g. a GPU has some level of hardware parallelism (Streaming Multiprocessor) and some levels of concurrency (warps and block of threads).
- Multidimensional range:** still experimental (as of January 2017), mapping a higher than 1D range of iteration.

```
// create the MDRangePolicy object  
using namespace Kokkos::Experimental;  
using range_type = MDRangePolicy< Rank<2>, Kokkos::IndexType<int> >;  
range_type range( {0,0}, {N0,N1} );
```

```
// use a special multidimensional parallel for launcher  
md_parallel_for(range, functor);
```



## Kokkos compute Kernels - parallel dispatch (8)

### Parallel dispatch - execution policy

- **Team policy:** for **hierarchical parallelism**

- threads team
- threads inside a team
- vector lanes

- ```
// Using default execution space and launching  
// a league with league_size teams with team_size threads each  
Kokkos::TeamPolicy <>  
policy( league_size , team_size );
```

equivalent to launching in CUDA a 1D grid of 1D blocks of threads.

Team scratch pad memory  $\iff$  CUDA shared memory



## Kokkos compute Kernels - parallel dispatch (9)

### Hierarchical parallelism (advanced)

- OpenMP: League of Teams of Threads
- Cuda: Grid of Blocks of Threads
- Experimental features: task parallelism  
see slides by C. Edwards at GTC2016 [2016-04-GTC-Kokkos-Task.pdf](#)



## Hands-On 2 : SAXPY

**Purpose:** The simplest computing kernel in Kokkos, importance of hwloc

- There 5 different versions
- **1. Serial : no Kokkos)**
- **2. OpenMP : no Kokkos)**
- 3. Kokkos-Lambda-CPU : Kokkos with lambda for threads dispatch
- **4. Kokkos-Lambda : Kokkos with lambda for threads dispatch and data buffer (Kokkos::View)**
- 5. Kokkos-Functor-CPU : Kokkos with functor for threads dispatch only
- **Saxpy serial (reference executable on Power8)**
  - `cd $HOME/kokkos-tutorial/kokkos-tutorials/1-Day-Tutorial/Exercises/01_AXPY/Serial`
  - `make KOKKOS_ARCH=Power8`
  - Alternatively, we could have modify Makefile and change SNB into Power8
- **Saxpy regular OpenMP (on Power8)**
  - `cd $HOME/kokkos-tutorial/kokkos-tutorials/1-Day-Tutorial/Exercises/01_AXPY/OpenMP`
  - Rebuild: `make KOKKOS_ARCH=Power8`; and observe performance

see also slides from SC2016, page 42(74).



## Hands-On 2 : SAXPY

### • Saxpy Kokkos OpenMP (on Power8) <sup>4</sup>

- `cd $HOME/kokkos-tutorial/kokkos-tutorials/1-Day-Tutorial/Exercises/01_AXPY/Kokkos-Lambda`
- Add 3 lines in `saxpy.cpp` right after Kokkos initialization

```
std::ostream msg;  
Kokkos::OpenMP::print_configuration( msg );  
std::cout << msg.str();
```

- `make KOKKOS_ARCH=Power8`
- Make sure all available CPU cores were used ( $1 \times 160 \times 1$ )
- Change the number of OpenMP threads created by kokkos, e.g. :  
`./saxpy.host -threads=20`
- Add again `KOKKOS_USE_TPLS="hwloc"` on the command line  
Rebuild and rerun, you should see that application uses **all the available numa domains**, and a strongly increased bandwidth usage !

---

<sup>4</sup> Make sure to use a very large data array; Power8 has very large cache memory. If you don't, this example will not measure memory bandwidth. Maximum bandwidth is 230 GB/s on a 2 socket P8. You should measure around 170 GB/s.



## Hands-On 2 : SAXPY


- **Saxpy CUDA (on Power8 + Nvidia K80/P100)**
  - `cd $HOME/kokkos-tutorial/kokkos-tutorials/1-Day-Tutorial/Exercises/01_AXPY/Kokkos-Lambda`
  - `module load cuda/8.0`
- Rebuild for K80, run on ouessant (front node):  
`make KOKKOS_DEVICES="Cuda,OpenMP"`  
`KOKKOS_ARCH="Kepler37,Power8" KOKKOS_USE_TPLS="hwloc"`
- Rebuild for P100, run on compute node using `submit_ouessant.sh` (should see a strong difference):  
`make KOKKOS_DEVICES="Cuda,OpenMP"`  
`KOKKOS_ARCH="Pascal60,Power8" KOKKOS_USE_TPLS="hwloc"`  
Please note that **maximun bandwith is 732 GB/s for Pascal P100**, you can retrieve this number by examining `deviceQuery` example in CUDA/SDK.



## Hands-On 3 : Mandelbrot set

- **Illustrate Functor class + 1D Kokkos::View + linearized index**
- the original **serial code** use 1D `std::vector<unsigned char>` data with linearized index, i.e.  $index = i + Nx * j$
- See **serial code** from `code/exercises/mandelbrot_kokkos/serial` (also read `main.cpp`)

```
for(int index=0; index<WIDTH*HEIGHT; ++index) {  
    int i,j;  
    index2coord(index,i,j,WIDTH,HEIGHT);  
    image[index]=mandelbrot(i,j);  
}
```





## Hands-On 3 : Mandelbrot set

### Proposed activity:

### refactor this computing loop into a C++ Kokkos functor class

- See [kokkos basic version](#) from `code/exercises/mandelbrot_kokkos/kokkos_basic` (already a bit refactored to ease the job)
- ❶ we added a file [kokkos\\_shared.h](#): `std::vector` replaced by a `Kokkos::View`
- ❷ **TODO:** fill TODOs in `mandelbrot.h` containing the definition of the c++ `mandelbrot kokkos functor`.  
**Notice:** the global constants have disappeared, they are now part of the functor context.
- ❸ **TODO:** refactor `main.cpp` (change the TODO)
  - Modify data allocation (from `std::vector` to `Kokkos::View`); we have now arrays: `image` and `imageHost` (mirror)
  - Copy back results from device to host.



## Hands-On 3 : Mandelbrot set

- The provided Makefile is designed to be used with kokkos environment from a modulefile
- Build the `kokkos_basic` version
- **OpenMP**
  - `module use /pwrwork/workshops/patc-201701/kokkos/modulefiles`
  - `module load kokkos/openmp_gnu485_dev`
  - `make`
- **Cuda**
  - `module use /pwrwork/workshops/patc-201701/kokkos/modulefiles`
  - `module load cuda/8.0 kokkos/cuda80_gnu485_dev_k80`
  - `make`
- **Compare performance** for a large Mandelbrot set  $8192 \times 8192$  : OpenMP versus Cuda



## Hands-On 3 : Mandelbrot set

- **Additional:** revisit this simple example using a **multidimensional range policy** to launch the Mandelbrot functor:

```
Kokkos::Experimental::MDRangePolicy< Kokkos::Experimental::Rank<2>  
                                     Kokkos::IndexType<int> >;
```

- **TODO:** fill TODOs in `mandelbrot.h` and `main.cpp` in directory `mandelbrot_kokkos/kokkos_mdrange`
- **This way avoids the use of linearized indexes.**



## Hands-On 3 : Mandelbrot set

- Pipelined version of Mandelbrot is not currently fully functional; it requires a small patch applied to Kokkos for cudaStreams;  
see <https://github.com/kokkos/kokkos/issues/532>
- Understand what is pipelined version of Mandelbrot see:  
<http://on-demand.gputechconf.com/gtc/2015/webinar/openacc-course/advanced-openacc-techniques.pdf>  
It basically consists in overlapping GPU computations with CPU/GPU memory transfert.
- See explanations given during training



## Hands-On 4 : Finite Difference / Stencil

- Illustrate the use of `2D Kokkos::View`



## MPI + Kokkos on Ouessant (1)

- Perform **distributed computing** on a cluster of **Power8 nodes** (4 GPU/node)
- **How to build application when KOKKOS\_DEVICE is Cuda ?**
  - Solution 1: Use `mpicxx` and pass env variable `OMPI_CXX=nvcc_wrapper`<sup>5</sup>
  - Solution 2: Use `nvcc_wrapper` as the compiler, but modify `CXX_FLAGS` / `LD_FLAGS` to add MPI specific flags.
- **How to make sure everything is ok regarding hardware affinity ?**  
**Cross-check at all possible level !** (so many ways to go wrong)
  - Use `mpirun --report-bindings` to cross-check afterwards how the job scheduler mapped the MPI task to core/host.
  - Use either `Kokkos::OpenMP::print_configuration / Kokkos::Cuda::print_configuration`
  - **Check MPI task - GPU binding is what you expect it to be in the application.**

```
int cudaDeviceId;  
cudaGetDevice(&cudaDeviceId);  
std::cout << "I'm MPI task #" << rank << " pinned to GPU #" << cudaDeviceId << "\n";
```

<sup>5</sup>Use `MPICH_CXX` if your MPI implementation is MPICH.

## MPI + Kokkos on Ouessant (2)

### Simple job script for using MPI+Kokkos/OpenMP

```
#!/bin/bash
#BSUB -x
#BSUB -J test_mpi_kokkos_openmp          # Job name
#BSUB -n 4                               # total number of MPI task
#BSUB -o test_mpi_kokkos_openmp.%J.out   # stdout filename
#BSUB -q compute                         # queue name
#BSUB -R "affinity[core(10):cpubind=core]" # affinity
#BSUB -R 'span[ptile=2]'                 # tile : number of MPI task/nod
#BSUB -W 00:05                           # maximum runtime
```

```
module load gcc/4.8/mpi/1.10
```

```
# number of OpenMP thread per MPI task
OMP_NUM_THREADS=20
```

```
# report bindings for cross-checking
```

```
mpirun --report-bindings -n ${LSB_DJOB_NUMPROC} ./test_mpi_kokkos.omp
```



## MPI + Kokkos on Ouessant (3)

### Simple job script for using MPI + Kokkos/Cuda

```
#!/bin/bash
#BSUB -x
#BSUB -J test_mpi_kokkos_cuda          # Job name
#BSUB -n 8                             # number of MPI tasks
#BSUB -o test_mpi_kokkos_cuda.%J.out   # stdout filename
#BSUB -q compute                       # queue name
#BSUB -R "affinity[core(5):cpubind=core]" # nb cores per MPI task
#BSUB -R "select[ngpus>0] rusage [ngpus_shared=1]" # activate GPU usage
#BSUB -W 00:05                         # max runtime
```

```
module load gcc/4.8/mpi/1.10 cuda/8.0
```

```
CUDA_VISIBLE_DEVICES=0,1,2,3
```

```
GPUS_PER_NODES=4
```

```
# Each mpi tasks are binded to a different GPU
```

```
mpirun --report-bindings -n ${LSB_DJOB_NUMPROC} ./test_mpi_kokkos.cuda --ndevices=$GPUS_PER_NODES
```

- This script requests 2 nodes, i.e.  $2 \times 4 = 8$  GPUs
- `core(5)`: just to be sure that each Power8 will receive 2 MPI tasks





## MPI + Kokkos on Ouessant (4) - Hands-On

### About LSF (job scheduler)

- Use code in `code/exercices/mpi_kokkos`; This application just reports bindings
- **Try to build this application against an installed version of Kokkos**, i.e. either OpenMP / Cuda
  - `module use /pwrwork/workshops/patc-201701/kokkos/modulefiles`
  - OpenMP: `module load kokkos/cuda80_gnu485_dev_k80`
  - Cuda: `module load cuda/8.0 kokkos/cuda80_gnu485_dev_k80`
  - `make`
  - This will build either `test_mpi_kokkos.omp` or `test_mpi_kokkos.cuda`
- Open and read `submit_ouessant_cpu.sh` / `submit_ouessant_gpu.sh`
- **Submit a job, read the output and check everything is what is expected**
- LSF commands to know:
  - **submit:** `bsub < submit_ouessant_cpu.sh`
  - **info/status:** `bjobs`
  - **cancel/kill:** `bkill`



## Hands-On 5 : Laplace solver with KOKKOS + MPI

**Slightly adapted/refactored from Nvidia's OpenACC exercise:**

[nvidia-advanced-openacc-course-sources](#)

We will use code from [code/exercises/laplace\\_kokkos](#), 4 different versions of the 2D Laplace solver:

- serial (no kokkos)
- **kokkos with 1D view** (linearized index)
- kokkos\_v2 with 2D views
- **kokkos\_mpi with MPI+CUDA and hwloc**



## Hands-On 6 - Reaction-Diffusion Fisher equation (1)

- **SETUP:** we will use git to download this miniApp code designed at [CSCS](#) for HPC teaching purpose.
  - `cd $HOME/patc_kokkos/code/miniapps/SummerSchool2016`
  - `git clone https://github.com/pkestene/SummerSchool2016.git`
  - `cd Summerschool2016; git checkout kokkos`
- **This material contains multiple versions** of a **Reaction-Diffusion PDE solver (Fisher equation)**. We will contribute 2 Kokkos versions of this solver.

$$\frac{\partial s}{\partial t} = D \left( \frac{\partial^2 s}{\partial x^2} + \frac{\partial^2 s}{\partial y^2} \right) + R s(1 - s) = 0$$



## Hands-On 6 - Reaction-Diffusion Fisher equation (2)

### 1 Explore/Read slides about the Fisher solver:

`$HOME/patc_kokkos/code/miniapps/SummerSchool2016/miniapp/kokkos/serial/miniapp.pdf`

- Explore the **serial** version of the Fisher solver.

### 2 These **Kokkos exercises** are routed to use the **modulefiles**:

- `module use /pwrwork/workshops/patc-201701/kokkos/modulefiles`
- `module load kokkos/openmp_gnu485_dev`
- `make`

### 3 **Kokkos version 1** / Exercice with KOKKOS\_LAMBDA / Already pre-filled, some TODOs

- Open and read file `miniapp/kokkos/cxx/readme.txt`
- Fill the TODO with Kokkos LAMBDA kernels



### 4 **Kokkos version 2** : already done

- The main difference between version 1 and 2 is how the c++ class `DataWarehouse` is designed
- Just build and compare performance with version 1, with Kokkos device OpenMP(Power8) and then Cuda



## Hands-On 7 : Euler equation solver

- Code location:  
`$HOME/patc_kokkos/code/miniapps/euler2d_kokkos_functor/`
- Build / run / mesure performance of a 2D Euler equation solver  
OpenMP/Cuda.
- See additionnal slides in source directory



## Using an installed Kokkos

- As you will surely **use multiple versions** of Kokkos (OpenMP, Cuda, ...), with/without Lambda, UVM, different compilers, etc ... it will be very useful to use some **modulefiles**.
- A **module environment** is not a tool specific to a super-computer, it can be used on a **Desktop/Laptop** to configure an execution environment.  
e.g. `sudo apt-get install environment-modules` (Debian/Ubuntu)
- **What is a modulefiles ?** A simple way to set env variables to ease the use of a given software package.
- You will find some examples modulefiles for Kokkos in  
`/pwrwork/workshops/patc-201701/kokkos/modulefiles/kokkos` you can easily adapt to your own platform.



## Using an installed Kokkos (2)

- A simple modulefiles for Kokkos should at minimum set variable `KOKKOS_PATH` pointing to the installed directory (the one which contains `Makefile.kokkos`)
- **How to use Kokkos modulefiles on your own machine ?** Just use the following:

```
# Assuming you placed the module file in  
# /somewhere_on_your_machince/modulefiles  
module use /somewhere_on_your_machince/modulefiles
```

```
# e.g. load Kokkos for GPU  
module load kokkos/cuda80_gnu485_dev_k80
```

- **How to use Kokkos modulefiles on ouessant ?** Just use the following:


```
# Assuming you placed the module file in  
# /somewhere_on_your_machince/modulefiles  
module use /pwrwork/workshops/patc-201701/kokkos/modulefiles  
# e.g. load Kokkos for GPU  
module load kokkos/cuda80_gnu485_dev_k80
```



## About Kokkos in Trilinos

- kokkos is originally a subpackage of trilinos (application framework for solving problems requiring parallel large distributed linear algebra solvers).
- **Kokkos is the performance portable layer**, to allow running Trilinos as efficiently as possible on multiple architectures.
- **Kokkos can be build independently from Trilinos** and used in other applications

Kokkos: Performance, Portability and Productivity



<https://github.com/kokkos>



## About Kokkos in Trilinos

- **Don't do the following on ouessant, your home is too small**, just keep the spirit to try on your own machine
- **Build a minimal featured Trilinos with Kokkos for GPU activated : Tpetra + kokkos + Cuda**
  - 1 **Example config platform:** Ubuntu 16.04 + openmpi + cuda 8.0  
compiler is gcc 5.4.0
  - 2 **Get Trilinos sources:**  

```
git clone https://github.com/trilinos/Trilinos.git; cd Trilinos;  
git checkout develop
```
  - 3 **CMake configuration script:** Use the provided configuration file `configure_tpetra_kokkos_cuda_nvcc_wrapper.sh` located in the provided archive (`doc/trilinos`)  
this script needs slight changes (`var OMPI_CXX` and `install prefix`)  
this script must be run in a build directory (not directly in trilinos sources).  
this config will build kokkos with unit tests and examples.
  - 4 **Build:** `make -j`; `make install`
  - 5 **Build a sample project.**



## Trilinos/Tpetra example project

- Directory `doc/trilinos/tpetra_example` contains a minimal example application for trilinos/tpetra. You just need to set env variable `TRILINOS_PATH` to install directory.



## Kokkos profiling interface (1)

- Kokkos provides by default a **profiling interface** through a **plugin mechanism**
- **Usage: profiling / monitoring / instrumenting**
- From an application point of view, there is nothing to do, just provide a plugin (shared library), e.g.

```
# define path to the plugin  
export KOKKOS_PROFILE_LIBRARY=/somewhere/kp_kernel_logger.so  
# run as usual Kokkos application
```

- Examples of Kokkos profile plugins can be found at  
<https://github.com/kokkos/kokkos-tools>



## Kokkos profiling interface (2)

- A Kokkos profile plugin must provide implementation for callback routines
  - `kokkosp_init_library`
  - `kokkosp_finalize_library`
- A Kokkos profile interface can provide implementation for callback routines specific to a type a parallel construct, e.g. `Kokkos::parallel_for`
  - `kokkosp_begin_parallel_for`
  - `kokkosp_end_parallel_for`

which are called every time application enters / exits this construct.

- see file `core/src/impl/Kokkos_Profiling_Interface.cpp` for a detailed list of possible callbacks.