

OpenMP Tasking In-depth

Christian Terboven

IT Center, RWTH Aachen University
terboven@itc.rwth-aachen.de

Loops with Tasks

- **Task generating construct: decompose a loop into chunks, create a task for each loop chunk**

```
#pragma omp taskloop [clause[[,] clause]...]  
{structured-for-loops}
```

- Where clause is one of:

- shared(list)
- private(list)
- firstprivate(list)
- lastprivate(list)
- default(sh | pr | fp | none)
- reduction(r-id: list)
- in_reduction(r-id: list)

Data Environment

- grainsize(grain-size)
- num_tasks(num-tasks)

Chunks/Grain

- if(scalar-expression)
- final(scalar-expression)
- mergeable

Cutoff Strategies

- untied
- priority(priority-value)

Scheduler (R/H)

- collapse(n)
- nogroup
- allocate([allocator:] list)

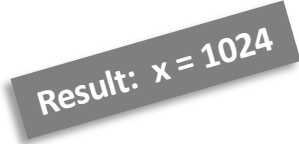
Miscellaneous

```
subroutine worksharing
  integer :: x
  integer :: i
  integer, parameter :: T = 16
  integer, parameter :: N = 1024

  x = 0
  !$omp parallel shared(x) num_threads(T)

  !$omp do
    do i = 1, N
      !$omp atomic
        x = x + 1
      !$omp end atomic
    end do
  !$omp end do

  !$omp end parallel
  write (*, '(A,I0)') 'x = ', x
end subroutine
```

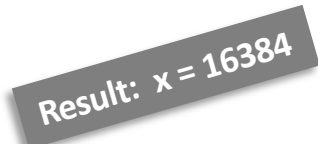


```
subroutine taskloop
  integer :: x
  integer :: i
  integer, parameter :: T = 16
  integer, parameter :: N = 1024

  x = 0
  !$omp parallel shared(x) num_threads(T)

  !$omp taskloop
    do i = 1, N
      !$omp atomic
        x = x + 1
      !$omp end atomic
    end do
  !$omp end taskloop

  !$omp end parallel
  write (*, '(A,I0)') 'x = ', x
end subroutine
```

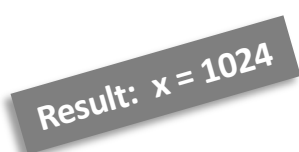


```
subroutine worksharing
  integer :: x
  integer :: i
  integer, parameter :: T = 16
  integer, parameter :: N = 1024

  x = 0
  !$omp parallel shared(x) num_threads(T)

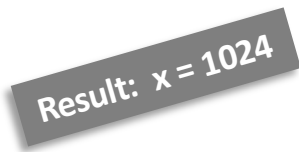
  !$omp do
    do i = 1, N
      !$omp atomic
        x = x + 1
      !$omp end atomic
    end do
  !$omp end do

  !$omp end parallel
  write (*, '(A,I0)') 'x = ', x
end subroutine
```



```
subroutine taskloop
  integer :: x
  integer :: i
  integer, parameter :: T = 16
  integer, parameter :: N = 1024

  x = 0
  !$omp parallel shared(x) num_threads(T)
  !$omp single
  !$omp taskloop
    do i = 1, N
      !$omp atomic
        x = x + 1
      !$omp end atomic
    end do
  !$omp end taskloop
  !$omp end single
  !$omp end parallel
  write (*, '(A,I0)') 'x = ', x
end subroutine
```



■ Clause: grainsize(grain-size)

- Chunks have at least grain-size iterations
- Chunks have maximum 2x grain-size

```
int TS = 4 * 1024;
#pragma omp taskloop grainsize(TS)
for ( i = 0; i<SIZE; i+=1) {
    A[i]=A[i]*B[i]*S;
}
```

■ Clause: num_tasks(num-tasks)

- Create num-tasks chunks
- Each chunk must have at least one iteration

```
int NT = 4 * omp_get_num_threads();
#pragma omp taskloop num_tasks(NT)
for ( i = 0; i<SIZE; i+=1) {
    A[i]=A[i]*B[i]*S;
}
```

- If none of previous clauses is present, the *number of chunks* and the *number of iterations per chunk* is implementation defined
- Additional considerations:
 - The order of the creation of the loop tasks is unspecified
 - Taskloop creates an implicit taskgroup region; **nogroup** → no implicit taskgroup region is created

Sudoku

■ Lets solve Sudoku puzzles with brute multi-core force

	6						8	11			15	14			16
15	11				16	14				12			6		
13		9	12					3	16	14		15	11	10	
2		16		11		15	10	1							
	15	11	10			16	2	13	8	9	12				
12	13			4	1	5	6	2	3					11	10
5		6	1	12		9		15	11	10	7	16			3
	2				10		11	6		5			13		9
10	7	15	11	16				12	13						6
9						1			2		16	10			11
1		4	6	9	13			7		11		3	16		
16	14			7		10	15	4	6	1				13	8
11	10		15				16	9	12	13			1	5	4
		12		1	4	6		16				11	10		
		5		8	12	13		10			11	2			14
3	16			10			7			6				12	

- (1) Search an empty field
- (2) Try all numbers:
 - (2 a) Check Sudoku
 - If invalid: skip
 - If valid:
Go to next field
- Wait for completion

■ Lets solve Sudoku puzzles with brute multi-core force

	6					8	11			15	14			16
15	11													
13		9	12											
2		16		11										
	15	11	10											
12	13			4									11	10
5		6	1	12		9		15	11	10	7	16		3
	2				10		11	6		5		13		9
10	7	15	11	16										
9														
1		4	6	9										
16	14			7		10	15	4	6	1			13	8
11	10		15				16	9	12	13			1	5
		12		1	4	6		16				11	10	
		5		8	12	13		10			11	2		14
3	16			10										

first call contained in a
`#pragma omp parallel`
`#pragma omp single`
such that one tasks starts the
execution of the algorithm

`#pragma omp task`
needs to work on a new copy
of the Sudoku board

`#pragma omp taskwait`
wait for all child tasks

- (1) Search an empty field
- (2) Try all numbers:
 - (2 a) Check Sudoku
 - If invalid: skip
 - If valid:
Go to next field
- Wait for completion

■ OpenMP parallel region creates a team of threads

```
#pragma omp parallel
{
    #pragma omp single
        solve_parallel(0, 0, sudoku2, false);
} // end omp parallel
```

→ Single construct: One thread enters the execution of `solve_parallel`

→ the other threads wait at the end of the `single` ...

→ ... and are ready to pick up threads „from the work queue“

■ Syntactic sugar (either you like it or you don't)

```
#pragma omp parallel sections
{
    solve_parallel(0, 0, sudoku2, false);
} // end omp parallel
```

■ The actual implementation

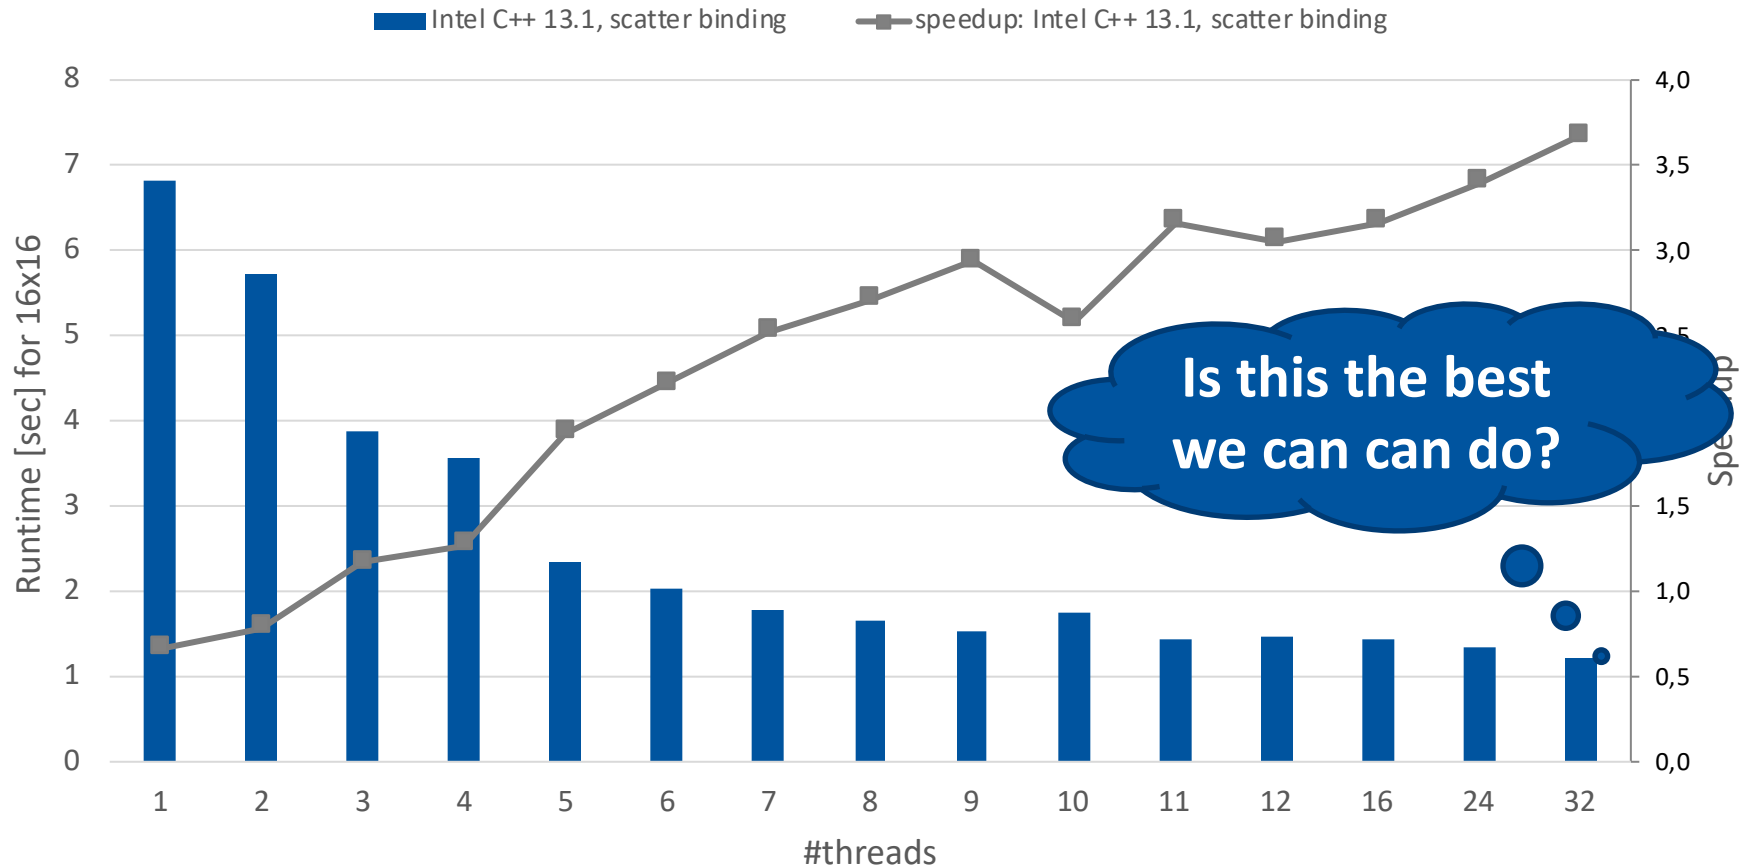
```
for (int i = 1; i <= sudoku->getFieldSize(); i++) {  
    if (!sudoku->check(x, y, i)) {  
#pragma omp task firstprivate(i,x,y,sudoku)  
    {  
        // create from copy constructor  
        CSudokuBoard new_sudoku(*sudoku)  
        new_sudoku.set(y, x, i);  
        if (solve_parallel(x+1, y, &new_sudoku)) {  
            new_sudoku.printBoard();  
        }  
    } // end omp task  
}
```

#pragma omp task
need to work on a new copy of
the Sudoku board

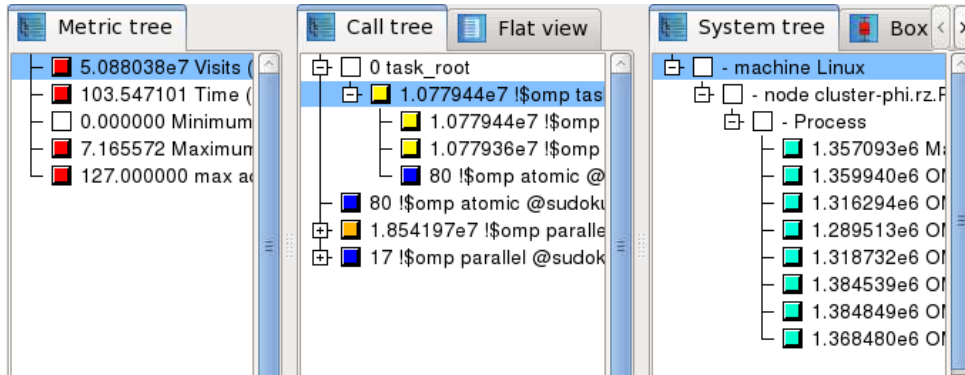
```
#pragma omp taskwait
```

#pragma omp taskwait
wait for all child tasks

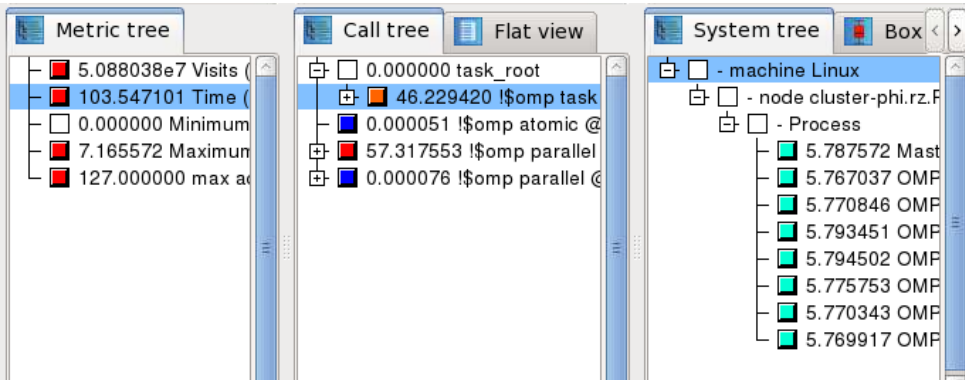
Sudoku on 2x Intel Xeon E5-2650 @2.0 GHz



Event-based profiling gives a good overview :



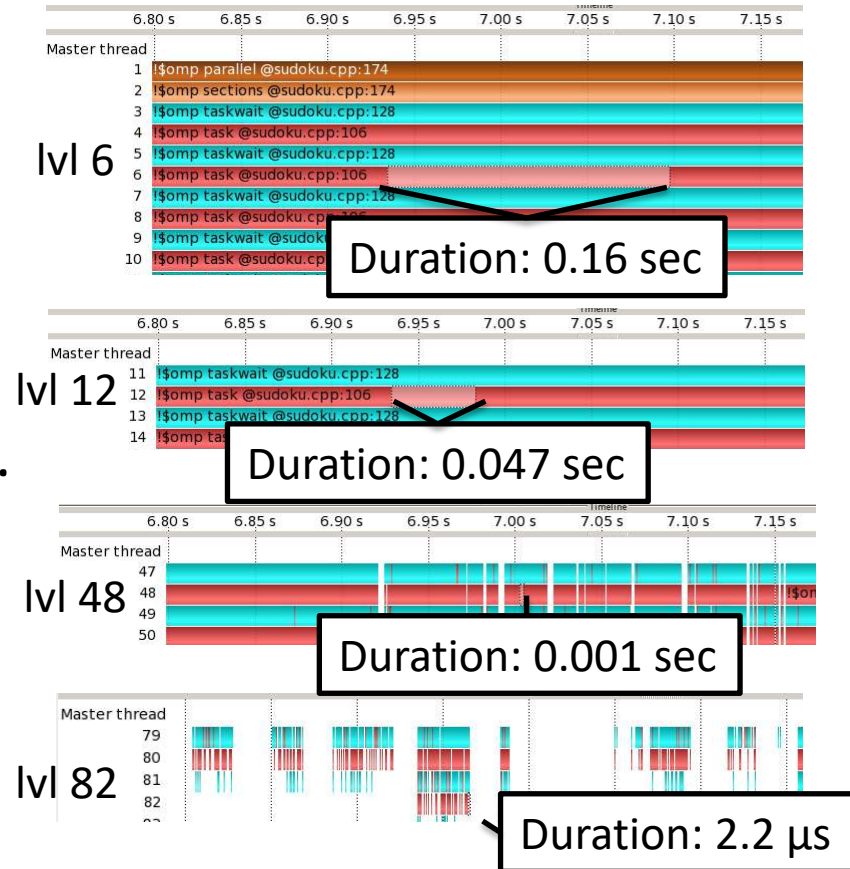
Every thread is executing ~1.3m tasks...



... in ~5.7 seconds.

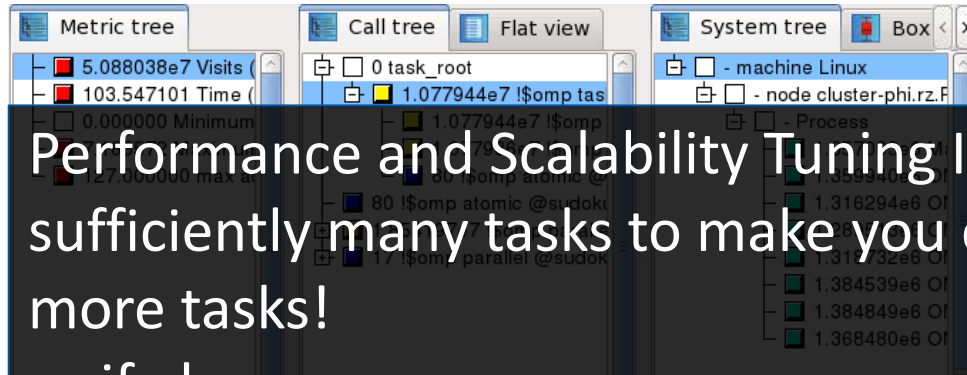
=> average duration of a task is ~4.4 μ s

Tracing gives more details:



Tasks get much smaller down the call-stack.

Event-based profiling gives a good overview :



Performance and Scalability Tuning Idea: If you have created sufficiently many tasks to make you cores busy, stop creating more tasks!

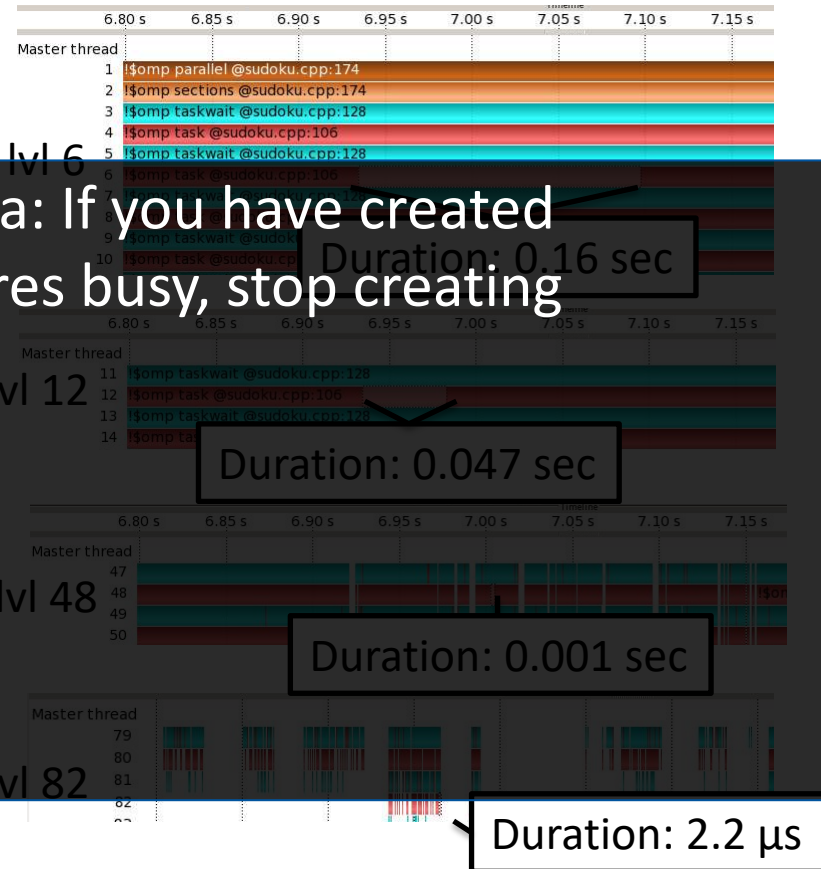
- if-clause
- final-clause, mergeable-clause
- natively in your program code

Example: stop recursion

... in ~5.7 seconds.

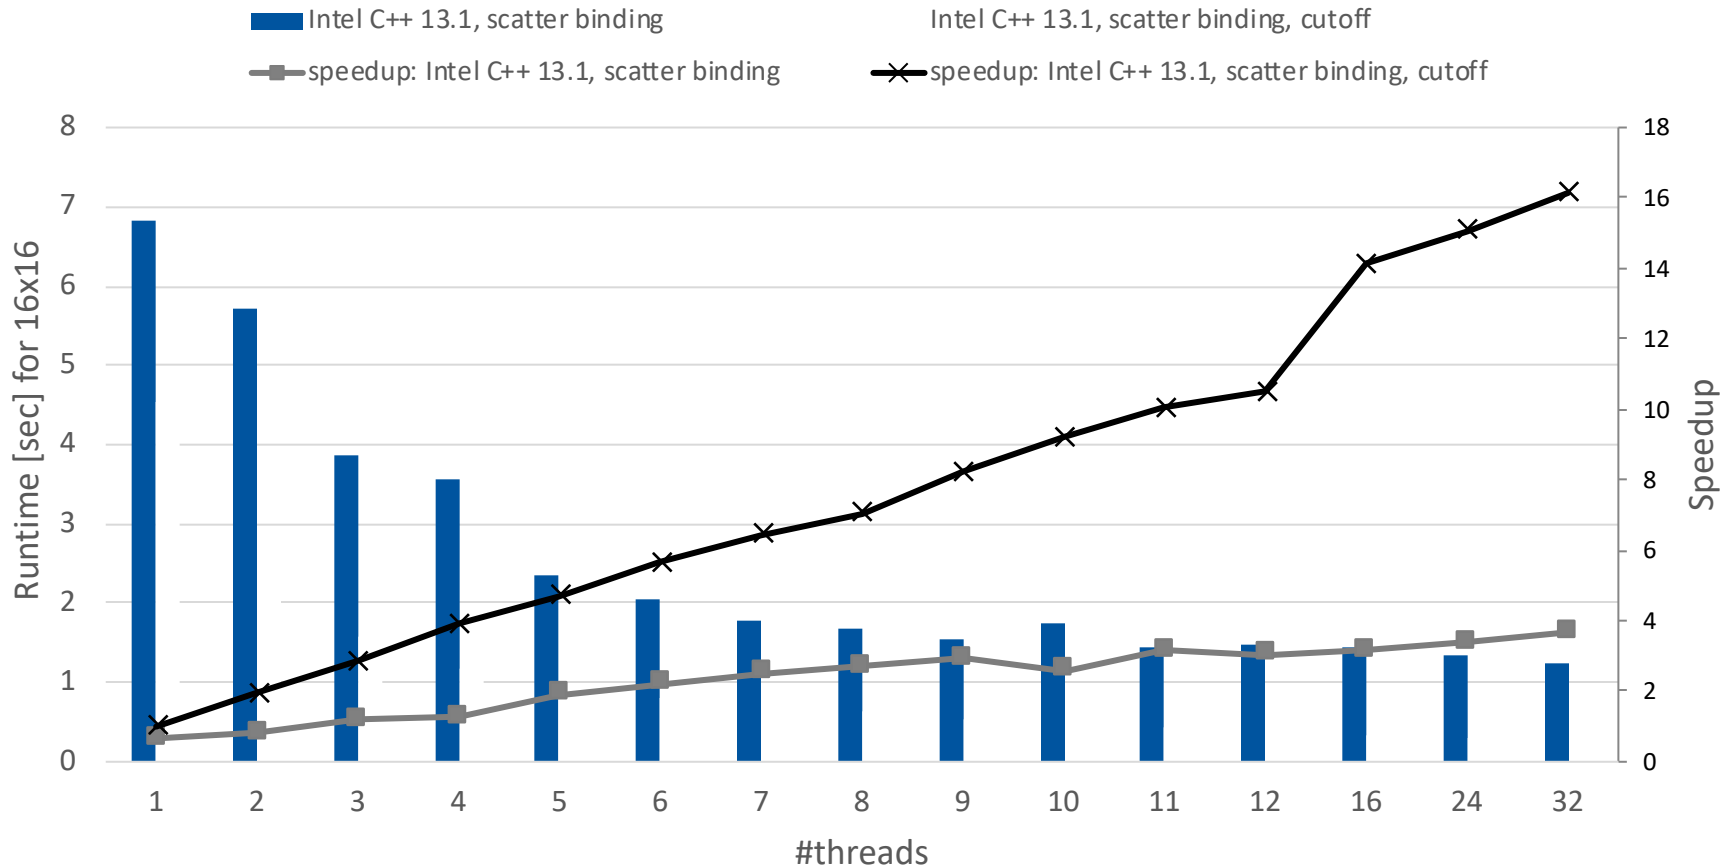
=> average duration of a task is ~4.4 μ s

Tracing gives more details:



Tasks get much smaller down the call-stack.

Sudoku on 2x Intel Xeon E5-2650 @2.0 GHz



Scheduling

- **Default: Tasks are *tied* to the thread that first executes them → not necessarily the creator. Scheduling constraints:**
 - Only the thread a task is tied to can execute it
 - A task can only be suspended at task scheduling points
 - Task creation, task finish, `taskwait`, `barrier`, `taskyield`
 - If task is not suspended in a barrier, executing thread can only switch to a direct descendant of all tasks tied to the thread
- **Tasks created with the `untied` clause are never tied**
 - Resume at task scheduling points possibly by different thread
 - ~~No scheduling restrictions, e.g., can be suspended at any point~~
 - But: More freedom to the implementation, e.g., load balancing

- **Problem:** Because untied tasks may migrate between threads at any point, thread-centric constructs can yield unexpected results
- **Remember when using untied tasks:**
 - Avoid `threadprivate` variable
 - Avoid any use of thread-ids (i.e., `omp_get_thread_num()`)
 - Be careful with `critical region` and *locks*
- **Simple Solution:**
 - Create a tied task region with

```
#pragma omp task if(0)
```

- The `taskyield` directive specifies that the current task can be suspended in favor of execution of a different task.

→ Hint to the runtime for optimization and/or deadlock prevention

C/C++

```
#pragma omp taskyield
```

Fortran

```
!$omp taskyield
```

taskyield Example (1/2)



```
#include <omp.h>

void something_useful();
void something_critical();

void foo(omp_lock_t * lock, int n)
{
    for(int i = 0; i < n; i++)
        #pragma omp task
        {
            something_useful();
            while( !omp_test_lock(lock) ) {
                #pragma omp taskyield
            }
            something_critical();
            omp_unset_lock(lock);
        }
}
```

taskyield Example (2/2)



```
#include <omp.h>

void something_useful();
void something_critical();

void foo(omp_lock_t * lock, int n)
{
    for(int i = 0; i < n; i++)
        #pragma omp task
        {
            something_useful();
            while( !omp_test_lock(lock) ) {
                #pragma omp taskyield
            }
            something_critical();
            omp_unset_lock(lock);
        }
}
```

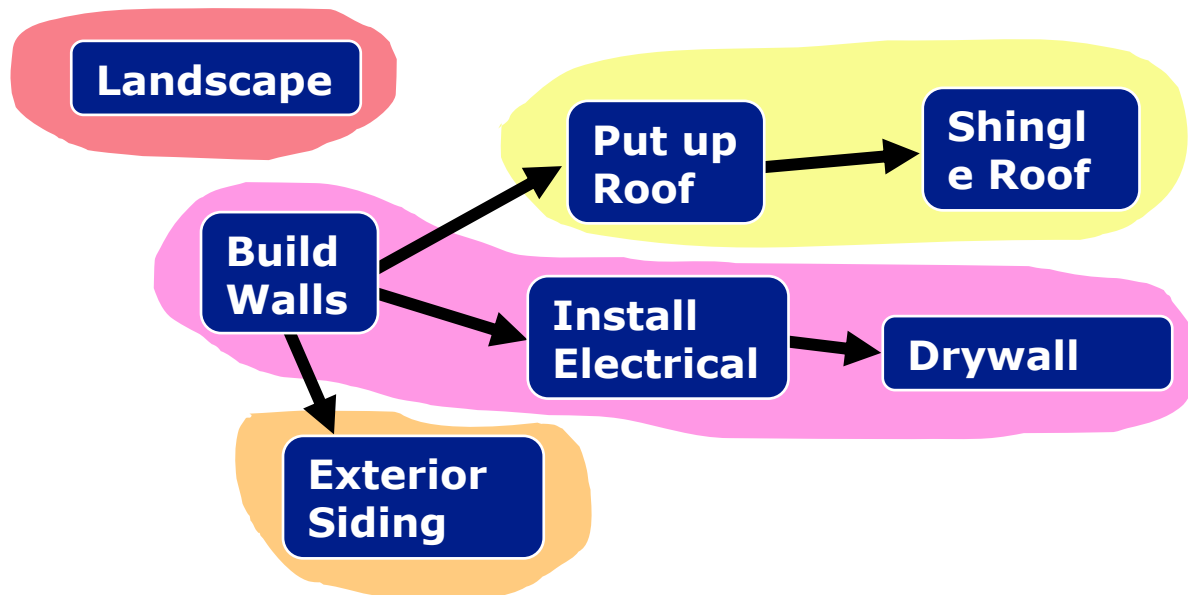
The waiting task may be suspended here and allow the executing thread to perform other work; may also avoid deadlock situations.

Tasks and Dependencies

Tasks and Dependencies



■ Catchy example: Building a house



- Task dependencies constrain execution order and times for tasks
- Fine-grained synchronization of tasks

```
int x = 0;
#pragma omp parallel
#pragma omp single
{
    ● #pragma omp task
      std::cout << x << std::endl;

    #pragma omp taskwait

    ● #pragma omp task
      x++;
}
```


Traditional task wait

```
int x = 0;
#pragma omp parallel
#pragma omp single
{
    ● #pragma omp task depend(in: x)
      std::cout << x << std::endl;




    ● #pragma omp task depend(inout: x)
      x++;
}
```

Dependencies

Task wait



Dependencies



Task's creation time

Task's execution time


More Complex Example: Cholesky Factorization



```
void cholesky(int ts, int nt, double* a[nt][nt])
{
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        potrf(a[k][k], ts, ts);

        // Triangular systems
        for (int i = k + 1; i < nt; i++)
            #pragma omp task
            trsm(a[k][k], a[k][i], ts, ts);
        #pragma omp taskwait

        // Update trailing matrix
        for (int i = k + 1; i < nt; i++)
            for (int j = k + 1; j < i; j++)
                #pragma omp task
                dgemm(a[k][i], a[k][j], a[j][i], ts, ts);
            #pragma omp task
            syrk(a[k][i], a[i][i], ts, ts);
        #pragma omp taskwait
    }
}
```




Traditional task wait

```
void cholesky(int ts, int nt, double* a[nt][nt]) {
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        #pragma omp task depend(inout: a[k][k])
        potrf(a[k][k], ts, ts);

        // Triangular systems
        for (int i = k + 1; i < nt; i++) {
            #pragma omp task depend(in: a[k][k])
            depend(inout: a[k][i])
            trsm(a[k][k], a[k][i], ts, ts);
        }

        // Update trailing matrix
        for (int i = k + 1; i < nt; i++) {
            for (int j = k + 1; j < i; j++) {
                #pragma omp task depend(inout: a[j][i])
                depend(in: a[k][i], a[k][j])
                dgemm(a[k][i], a[k][j], a[j][i], ts, ts);
            }
            #pragma omp task depend(inout: a[i][i])
            depend(in: a[k][i])
            syrk(a[k][i], a[i][i], ts, ts);
        }
    }
}
```



Dependencies

Questions?